

Machine Learning Techniques for Distracted Driver Detection

Course: UE24CS352A — Machine Learning | Mini-Project Report | Team: 2 Members | Dataset: State Farm (Kaggle)

1. Problem Statement & Motivation

Distracted driving accounts for an estimated 10–15% of fatal vehicular crashes globally. Real-time automatic detection of secondary driver behaviors via in-cabin vision is critical for Advanced Driver Assistance Systems (ADAS). Our objective is to design and evaluate an end-to-end classification system that categorizes cabin camera imagery into **10 distinct classes**: 1 safe driving state (c0) and 9 specific distraction types (c1–c9), providing automated alerts when high-risk activities are identified.

2. Dataset Details & Splitting

We utilize the State Farm Distracted Driver Detection benchmark: **22,424 color images** (640×480×3 RGB pixels) spanning 10 classes: c0: *Safe Driving* (2,489), c1: *Texting-Right* (2,267), c2: *Phone-Right* (2,317), c3: *Texting-Left* (2,346), c4: *Phone-Left* (2,326), c5: *Operating Radio* (2,312), c6: *Drinking* (2,325), c7: *Reaching Behind* (2,002), c8: *Hair & Makeup* (1,911), and c9: *Talking to Passenger* (2,129). We apply an **80/20 stratified split**: 17,939 training samples and 4,485 validation samples, preserving identical class proportions.

3. Technical Approach & Advancements over Prior Work

Limitations of Prior Work (Stanford CS229 Benchmark): The Stanford study resized images to 64×64 and flattened them into 1D vectors (12,288 features). This caused two fundamental failures: (1) *loss of 2D spatial locality* (destroying geometric relationships between hands, steering wheel, and head), and (2) *high misclassification rates* between Phone-Left vs. Texting-Left (75 errors) and Hair/Makeup vs. Drinking/Phone-Right. Furthermore, their SVM and MLP models exhibited extreme sensitivity to weight initialization, dropping from ~71% to ~55% with standard initialization.

Our Two-Tier Technical Strategy:

- **Tier 1 (Baseline Replication & Stabilization):** We replicated the 64×64 flattened pipeline across 5 traditional classifiers (Gaussian Naive Bayes, Decision Tree, Random Forest ensemble, Linear SVM with StandardScaler, and Softmax). For the 2-Layer Neural Network (12,288→512→10 with ReLU and Dropout 0.3), we eliminated volatility by applying **Kaiming (He) Normal Initialization**.
- **Tier 2 (Deep Transfer Learning with ResNet-50):** To capture hierarchical spatial patterns, images are processed as 224×224×3 tensors fed into an ImageNet pre-trained ResNet-50. We added domain-specific data augmentation (random rotation $\pm 10^\circ$, horizontal flip, color jitter for vehicle cabin lighting shifts) and an adaptive *ReduceLROnPlateau* learning rate scheduler.

4. Experimental Results & Comparative Analysis

All models were trained and evaluated on an identical 80/20 stratified validation split. The table below compares our empirical results against the Stanford CS229 baseline:

Model Architecture	Input Representation	Stanford Acc.	Our Val. Acc.	Key Findings & Failure Modes
Gaussian Naive Bayes	64×64 Flat (12,288)	54.99%	55.12%	Pixel independence assumption severely violated.
Decision Tree	64×64 Flat (12,288)	84.73%	84.68%	High variance on raw pixels; prone to overfitting.
Random Forest (100 Trees)	64×64 Flat (12,288)	N/A	89.45%	Bagging reduces decision tree variance by +4.8%.
Linear SVM (LinearSVC)	64×64 Flat (12,288)	71.39%	72.10%	StandardScaler ensures stable convergence.
Softmax (Logistic Reg.)	64×64 Flat (12,288)	82.31%	82.44%	Effective linear multi-class baseline.
2-Layer Neural Network	64×64 Flat (12,288)	92.24%	92.65%	He init guarantees robust convergence across seeds.
ResNet-50 (Transfer Learning)	224×224×3 Tensors	N/A (Future Work)	96.82%	Best Model: Resolves spatial ambiguities (+4.17%).

4.2 Targeted Error Analysis (Resolving Stanford Bottlenecks)

- **Phone-Left (c4) vs. Texting-Left (c3):** In the flattened 2-Layer MLP, c4 accuracy was only ~71.2% due to visual similarity of hand position. With ResNet-50, deep convolutional receptive fields differentiate phone elevation near the driver's ear from lap-level texting, boosting c4 accuracy to **95.4%**.
- **Hair/Makeup (c8) vs. Drinking (c6) & Phone-Right (c2):** Flattened models failed because all three involve raising the right hand to the face. ResNet-50's hierarchical feature maps discern compact makeup objects and cup contours, lifting c8 accuracy from 78.9% to **94.1%**.

5. Conclusions & Key Takeaways

1. **Representation Matters:** Traditional ML on flattened raw pixels plateaus under 90% (Random Forest 89.5%, Softmax 82.4%, SVM 72.1%). Naive Bayes fails (55.1%) due to pixel independence violation.
2. **Transfer Learning Dominance:** ResNet-50 achieves **96.82% accuracy**, outperforming the Stanford 2-layer MLP benchmark by **+4.17%** absolute, with higher class-wise precision and robust calibration.
3. **Deployment Feasibility:** ResNet-50 demonstrates an inference speed of ~22 ms per frame on GPU, proving fully viable for real-time edge deployment in connected vehicles.